

# COMS30035, Machine learning: Regression and Classification Trees

James Cussens

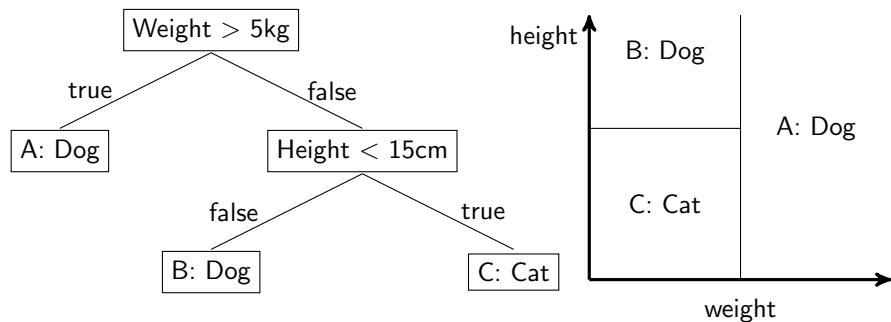
School of Computer Science  
University of Bristol

3rd September 2024

# Acknowledgement

- ▶ These slides are adapted from ones originally created by Edwin Simpson.

# Decision Trees



# Decision Trees as Partitioning Input Space

- ▶ One model is responsible for assigning a decision for each region of input space;
- ▶ The correct model for an input  $\mathbf{x}$  is chosen by traversing the binary decision tree, following the path from the top to a leaf.
- ▶ Leaf node is responsible for assigning a decision, such as a:
  - ▶ Class label;
  - ▶ Probability distribution over class labels;
  - ▶ Scalar value (for regression tasks).

# Learning the Tree Structure

- ▶ Which input variable to use at each node?
- ▶ What threshold to set for the split at each node?
- ▶ Classification and Regression Trees (CART): one of many possible learning algorithms
- ▶ Objective: greedily minimise the error
  - ▶ Regression: sum-of-squares
  - ▶ Classification: cross-entropy as used in neural networks or Gini impurity

# Learning the Tree Structure

- ▶ Which input variable to use at each node?
- ▶ What threshold to set for the split at each node?
- ▶ Classification and Regression Trees (CART): one of many possible learning algorithms
- ▶ Objective: greedily minimise the error
  - ▶ Regression: sum-of-squares
  - ▶ Classification: cross-entropy as used in neural networks or Gini impurity

# Learning the Tree Structure

- ▶ Which input variable to use at each node?
- ▶ What threshold to set for the split at each node?
- ▶ Classification and Regression Trees (CART): one of many possible learning algorithms
- ▶ Objective: greedily minimise the error
  - ▶ Regression: sum-of-squares
  - ▶ Classification: cross-entropy as used in neural networks or Gini impurity

# Learning the Tree Structure

- ▶ Which input variable to use at each node?
- ▶ What threshold to set for the split at each node?
- ▶ Classification and Regression Trees (CART): one of many possible learning algorithms
- ▶ Objective: greedily minimise the error
  - ▶ Regression: sum-of-squares
  - ▶ Classification: cross-entropy as used in neural networks or Gini impurity



# Learning the Tree Structure

- ▶ Number of possible solutions grows combinatorially with the number of input variables
- ▶ Greedy algorithm: add nodes one-at-a-time, choosing the best split at each point
  1. Start from the root node
  2. Run *exhaustive search* over each possible variable and threshold for a new node. For each variable and threshold:
    - ▶ Compute average of the target variable for each leaf of the proposed node
    - ▶ Compute the error if we stop adding nodes here
  3. Choose the variable & threshold that minimise the error
  4. Add a new node for the chosen variable and threshold.
  5. Repeat step 2 until there are only  $n$  data points associated with each leaf node.
  6. Prune back the tree to remove branches that do not reduce error by more than a small tolerance value,  $\epsilon$ .

# Learning the Tree Structure

- ▶ Number of possible solutions grows combinatorially with the number of input variables
- ▶ Greedy algorithm: add nodes one-at-a-time, choosing the best split at each point
  1. Start from the root node
  2. Run *exhaustive search* over each possible variable and threshold for a new node. For each variable and threshold:
    - ▶ Compute average of the target variable for each leaf of the proposed node
    - ▶ Compute the error if we stop adding nodes here
  3. Choose the variable & threshold that minimise the error
  4. Add a new node for the chosen variable and threshold.
  5. Repeat step 2 until there are only  $n$  data points associated with each leaf node.
  6. Prune back the tree to remove branches that do not reduce error by more than a small tolerance value,  $\epsilon$ .

# Learning the Tree Structure

- ▶ Number of possible solutions grows combinatorially with the number of input variables
- ▶ Greedy algorithm: add nodes one-at-a-time, choosing the best split at each point
  1. Start from the root node
  2. Run *exhaustive search* over each possible variable and threshold for a new node. For each variable and threshold:
    - ▶ Compute average of the target variable for each leaf of the proposed node
    - ▶ Compute the error if we stop adding nodes here
  3. Choose the variable & threshold that minimise the error
  4. Add a new node for the chosen variable and threshold.
  5. Repeat step 2 until there are only  $n$  data points associated with each leaf node.
  6. Prune back the tree to remove branches that do not reduce error by more than a small tolerance value,  $\epsilon$ .

# Learning the Tree Structure

- ▶ Number of possible solutions grows combinatorially with the number of input variables
- ▶ Greedy algorithm: add nodes one-at-a-time, choosing the best split at each point
  1. Start from the root node
  2. Run *exhaustive search* over each possible variable and threshold for a new node. For each variable and threshold:
    - ▶ Compute average of the target variable for each leaf of the proposed node
    - ▶ Compute the error if we stop adding nodes here
  3. Choose the variable & threshold that minimise the error
  4. Add a new node for the chosen variable and threshold.
  5. Repeat step 2 until there are only  $n$  data points associated with each leaf node.
  6. Prune back the tree to remove branches that do not reduce error by more than a small tolerance value,  $\epsilon$ .

# Learning the Tree Structure

- ▶ Number of possible solutions grows combinatorially with the number of input variables
- ▶ Greedy algorithm: add nodes one-at-a-time, choosing the best split at each point
  1. Start from the root node
  2. Run *exhaustive search* over each possible variable and threshold for a new node. For each variable and threshold:
    - ▶ Compute average of the target variable for each leaf of the proposed node
    - ▶ Compute the error if we stop adding nodes here
  3. Choose the variable & threshold that minimise the error
  4. Add a new node for the chosen variable and threshold.
  5. Repeat step 2 until there are only  $n$  data points associated with each leaf node.
  6. Prune back the tree to remove branches that do not reduce error by more than a small tolerance value,  $\epsilon$ .

# Learning the Tree Structure

- ▶ Number of possible solutions grows combinatorially with the number of input variables
- ▶ Greedy algorithm: add nodes one-at-a-time, choosing the best split at each point
  1. Start from the root node
  2. Run *exhaustive search* over each possible variable and threshold for a new node. For each variable and threshold:
    - ▶ Compute average of the target variable for each leaf of the proposed node
    - ▶ Compute the error if we stop adding nodes here
  3. Choose the variable & threshold that minimise the error
  4. Add a new node for the chosen variable and threshold.
  5. Repeat step 2 until there are only  $n$  data points associated with each leaf node.
  6. Prune back the tree to remove branches that do not reduce error by more than a small tolerance value,  $\epsilon$ .

# Learning the Tree Structure

- ▶ Number of possible solutions grows combinatorially with the number of input variables
- ▶ Greedy algorithm: add nodes one-at-a-time, choosing the best split at each point
  1. Start from the root node
  2. Run *exhaustive search* over each possible variable and threshold for a new node. For each variable and threshold:
    - ▶ Compute average of the target variable for each leaf of the proposed node
    - ▶ Compute the error if we stop adding nodes here
  3. Choose the variable & threshold that minimise the error
  4. Add a new node for the chosen variable and threshold.
  5. Repeat step 2 until there are only  $n$  data points associated with each leaf node.
  6. Prune back the tree to remove branches that do not reduce error by more than a small tolerance value,  $\epsilon$ .

# Pruning

- ▶ Balance residual training-set error against model complexity
- ▶ Start with a tree  $T_0$
- ▶ Consider pruning each node in  $T_0$  by combining the branches to obtain tree  $T$
- ▶ Compute a criterion  $C(T) = \sum_{\tau=1}^{|T|} e_{\tau}(T) + \lambda|T|$
- ▶ If  $C(T) \leq C(T_0)$  keep the pruned tree, else reinstate the pruned node.



# Pruning

- ▶ Balance residual training-set error against model complexity
- ▶ Start with a tree  $T_0$
- ▶ Consider pruning each node in  $T_0$  by combining the branches to obtain tree  $T$
- ▶ Compute a criterion  $C(T) = \sum_{\tau=1}^{|T|} e_{\tau}(T) + \lambda|T|$
- ▶ If  $C(T) \leq C(T_0)$  keep the pruned tree, else reinstate the pruned node.

# Pruning

- ▶ Balance residual training-set error against model complexity
- ▶ Start with a tree  $T_0$
- ▶ Consider pruning each node in  $T_0$  by combining the branches to obtain tree  $T$
- ▶ Compute a criterion  $C(T) = \sum_{\tau=1}^{|T|} e_{\tau}(T) + \lambda|T|$
- ▶ If  $C(T) \leq C(T_0)$  keep the pruned tree, else reinstate the pruned node.

# Pruning

- ▶ Balance residual training-set error against model complexity
- ▶ Start with a tree  $T_0$
- ▶ Consider pruning each node in  $T_0$  by combining the branches to obtain tree  $T$
- ▶ Compute a criterion  $C(T) = \sum_{\tau=1}^{|T|} e_{\tau}(T) + \lambda|T|$
- ▶ If  $C(T) \leq C(T_0)$  keep the pruned tree, else reinstate the pruned node.

# Interpretability

- ▶ The sequence of decisions is often easier to interpret than other methods (think of neural networks);
- ▶ However, sometimes small changes to the dataset cause big changes to the tree;
- ▶ If the optimal decision boundary is not aligned with the axes of an input variable, we need a lot of splits.

